



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

Crossplane vs Kubebuilder

paradigms, abstraction, complexity, extensibility and performance comparison with a common
use case implementation

Infrastructures for Cloud Computing and Big Data M
Lorenzo De Luca
0001240461

Kubernetes Operators



A Kubernetes Operator extends the K8s API with a Custom Resource (CR) and a controller that continuously reconciles actual state toward desired state.

Building operators can be approached at different levels of abstraction:

Low-level: the developer writes all the controller logic

High-level: the developer describe what he wants and let a platform generate the controller logic

Both approaches solve the same problem of automating resource lifecycle but target different users and different tradeoffs.



Research Objective

Compare Crossplane and Kubebuilder to compare the two paradigms for building Kubernetes Operators.

Will be evaluated:

- Architectural model and abstraction level
- Development complexity and extensibility
- Internal optimizations adopted by each framework
- Resource consumption


does the convenience of the extra abstraction layer of Crossplane cost more in resources/latency than it saves in development effort, compared to a hand-written Kubebuilder controller?




Crossplane



- Universal Control Plane
- Declarative composition model
- Strong for composition and platform engineering
- Logic in YAML + patch/transform functions
- Strength: consistency and reuse
- Limitation: less flexible for complex imperative logic



Kubebuilder



- Low-level Go SDK
- Imperative control model
- Strong for writing specialized controllers from scratch.
- Logic in Go
- Strength: maximum flexibility
- Limitation: higher coding and maintenance effort



Features comparison

Aspect	Crossplane	Kubebuilder
Abstraction	High(declarative)	Low(imperative)
Architecture	Multi-layer	Single controller
Development effort	Lower	Higher
Extensibility	Providers, Functions, patching	Arbitrary code, any library
Flexibility	Limited	Very high
Reusability	High	Low
Performance	Higher overhead	More efficient
Debugging	Harder	Easier(centralized)
Best	Kubebuilder	Less layers, more control

Optimizations


Crossplane:

- Provider-level connection reuse/pooling toward external APIs
- Leader election prevents duplicate reconciliation in multi-replica setups
- Multiple controllers run concurrently (core + each provider + functions)



Kubebuilder

- Informer cache: reconciler reads from local cache, not the API server
- Work queue with exponential backoff on repeated failures
- Managed Resource state cached; reconciliation skipped when observed state already matches desired state



So... Which one is better? It depends...

Scenario	Better choice	Reason
Standard infrastructure composition	Crossplane	Declarative and reusable
Internal Developer Platform	Crossplane	User-friendly claims and abstractions
Multi-resource self-service APIs	Crossplane	Composition model fits well
Complex conditional logic	Kubebuilder	Full control in Go
Custom external API integration	Kubebuilder	Any Go library can be used
Fine-grained lifecycle control	Kubebuilder	Reconciliation is fully programmable
Provisioning latency	Kubebuilder	Fewer controller layers
Debuggability	Kubebuilder	Less layers, more control

Kafka Topic Self-Service

<https://github.com/lorenzodeluca/kafka-operator-clash>

Common goal





The goal for the implemented operators is to create a new Kafka topic. Instead of manually creating all required resources, the user declares a single object called *DataStream*.

When a *DataStream* is created, the operator must automatically provision and reconcile a ready-to-use data pipeline: a Kafka topic where data can flow, a configuration file with the connection details, and a small consumer application that reads the connection details generated by the operator.

DataStream example:

```
1  apiVersion: messaging.lorenzodeluca.it/v1alpha1
2  kind: DataStream
3  metadata:
4    name: lorenzo-sandbox-stream
5    namespace: default
6  spec:
7    topicName: "user-telemetry-data"
8    partitions: 2
9    replicationFactor: 1
```

Crossplane implementation




api/datastream_xrd.yaml:

```
1  apiVersion: apiextensions.crossplane.io/v1
2  kind: CompositeResourceDefinition
3  metadata:
4    name: xdatastreams.messaging.lorenzodeluca.it
5  spec:
6    group: messaging.lorenzodeluca.it
7    names:
8      kind: XDataStream
9      plural: xdatastreams
10   claimNames:
11     kind: DataStream
12     plural: datastreams
13   versions:
14     - name: v1alpha1
15       served: true
16       referenceable: true
17       schema:
18         openAPIV3Schema:
19           type: object
20           properties:
21             spec:
22               type: object
23               properties:
24                 topicName:
25                   type: string
26                 partitions:
27                   type: integer
28                   default: 1
29                 replicationFactor:
30                   type: integer
31                   default: 1
32               required:
33                 - topicName
```

compositions/datastream_composition.yaml

```
- step: patch-and-transform
  functionRef:
    name: function-patch-and-transform
  input:
    apiVersion: pt.fn.crossplane.io/v1beta1
    kind: Resources
    resources:
      - name: kafka-topic
        base:
          apiVersion: topic.kafka.crossplane.io/v1alpha1
          kind: Topic
          spec:
            providerConfigRef:
              name: kafka-default
            forProvider:
              partitions: 1
              replicationFactor: 1
        patches:
          - type: FromCompositeFieldPath
            fromFieldPath: spec.topicName
            toFieldPath: metadata.name
          - type: FromCompositeFieldPath
            fromFieldPath: spec.partitions
            toFieldPath: spec.forProvider.partitions
          - type: FromCompositeFieldPath
            fromFieldPath: spec.replicationFactor
            toFieldPath: spec.forProvider.replicationFactor
```

Kubebuilder implementation



datastream_types.go:

```
1 package v1alpha1
2
3 import (
4     metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"
5     "k8s.io/apimachinery/pkg/runtime"
6 )
7
8 // DataStreamSpec defines the desired state of DataStream
9 type DataStreamSpec struct {
10     // TopicName is the name of the Kafka topic to manage.
11     // +kubebuilder:validation:MinLength=1
12     TopicName string `json:"topicName"`
13
14     // Partitions defines the total partition count for the topic.
15     // +kubebuilder:default:=1
16     // +kubebuilder:validation:Minimum=1
17     Partitions int32 `json:"partitions,omitEmpty"`
18
19     // ReplicationFactor defines the replication factor.
20     // +kubebuilder:default:=1
21     // +kubebuilder:validation:Minimum=1
22     ReplicationFactor int16 `json:"replicationFactor,omitEmpty"`
23 }
24
```

datastream_controller.go:

```
153 func createKafkaTopic(broker, topic string, partitions int32, reps int16) error {
154     conn, err := kafka.Dial("tcp", broker)
155     if err != nil {
156         return fmt.Errorf("dial broker failed: %w", err)
157     }
158     defer conn.Close()
159
160     controller, err := conn.Controller()
161     if err != nil {
162         return fmt.Errorf("cannot get kafka controller: %w", err)
163     }
164
165     controllerAddr := net.JoinHostPort(controller.Host, strconv.Itoa(controller.Port))
166     controllerConn, err := kafka.Dial("tcp", controllerAddr)
167     if err != nil {
168         controllerConn, err = kafka.Dial("tcp", broker)
169         if err != nil {
170             return fmt.Errorf("cannot dial kafka controller (%s) nor broker fallback (%s): %w", controllerAddr, broker, err)
171         }
172     }
173     defer controllerConn.Close()
174
175     err = controllerConn.CreateTopics(kafka.TopicConfig{
176         Topic:          topic,
177         NumPartitions:  int(partitions),
178         ReplicationFactor: int(reps),
179     })
180     if err != nil {
181         l := strings.ToLower(err.Error())
182         if strings.Contains(l, "already exists") || strings.Contains(l, "topic already exists") {
183             return nil
184         }
185         return err
186     }
187     return nil
188 }
```

Performance results


Metric	Crossplane	Kubebuilder
Lines of code written*	165	270
Edited files	7	2
Avg time to ready state (seconds)**	6	2

- Excluding generated code and provider installation/configuration manifests.

** The measured time is the elapsed time between applying the DataStream resource and observing the expected ready state.


Environment:

- Debian vps with 8 cores cpu, 8Gb ram, 75Gb ssd
- Local k3s cluster
- Single Kafka broker
- One DataStream resource per run
- Controllers/providers already running




Resource consumption

Average CPU usage vs DataStream objects batch size




Average RAM usage vs DataStream objects batch size



X-axis: batch size, the number of DataStream resources created together


Y-axis: average CPU rate (millicores)/RAM usage each component sustained while it was actively reconciling that batch

*Crossplane measurements includes the full control plane (core + providers + functions).




Resource consumption Crossplane

average CPU usage(millicores) vs batch size




average RAM usage(Mib) vs DataStream objects batch



X-axis: batch size, the number of DataStream resources created together (N)

Y-axis: average CPU rate (millicores)/RAM usage(MiB) each component sustained while it was actively reconciling that batch



REFERENCES

- The Kubebuilder Book, <https://book.kubebuilder.io>
- Crossplane documentation, <https://docs.crossplane.io/v2.3>
- kafka-go, <https://github.com/segmentio/kafka-go>



Questions?

Thank you for your attention

